

NYU, Tandon School of Engineering

October' 2024

CS-GY 6083

Principles of Database Systems
Section A, Fall 2024

Homework #3

Submitted by:

Khwaab Thareja | N15911999 | Kt3180

Guided By: Prof Phyllis Frankl

Use the University DB, unless otherwise noted. Write SQL queries (unless otherwise noted) for each of the following:

1. Find IDs and names of students who got a B in CS-101 and an A in CS-319

Answer 1

```
Select distinct s.ID,name
from student s
inner join takes t on s.ID=t.ID
where t.course_id='CS-319' and Grade='A'
and s.id in
( Select s.ID from student s
inner join takes t on s.ID=T.ID
and t.course_id='CS-101' and Grade='B')
```

2. A. Create a table Gradepoint(grade,points) to associate letter grades with points and fill it with the appropriate values ('A', 4.0), ('A-', 3.7), etc.

B. Add a foreign key constraint to another table, referencing this Gradepoint table. Note which table, write the constraint, including “on delete” and “on update” clauses and briefly justify your choices for those clauses.

C. Define a VIEW GradePointAvg(ID, GPA) that lists each student’s ID and gradepoint average.

Note: If all courses had the same number of credits, you could compute gradepoint averages with a query involving a natural join of takes and gradepoint, along with the AVG aggregation operator, grouping by ids. However, that solution doesn’t take account of different courses having different numbers of credits. Assume that all graded courses are included, even if a student repeats the same course.

Answer 2.A

Drop table if exists GradePoint;

create table GradePoint

 (grade varchar(2),

 point float(4),

 primary key (grade)

);

insert into GradePoint values ('A', 4.0);

insert into GradePoint values ('A-', 3.7);

insert into GradePoint values ('B+', 3.3);

insert into GradePoint values ('B', 3.0);

insert into GradePoint values ('B-', 2.7);

insert into GradePoint values ('C+', 2.3);

insert into GradePoint values ('C', 2.0);

insert into GradePoint values ('C-', 1.7);

insert into GradePoint values ('D+', 1.3);

insert into GradePoint values ('D', 1.0);

insert into GradePoint values ('F', 0);

Answer 2B

```
ALTER TABLE Takes
```

```
ADD CONSTRAINT fk_grade
```

```
FOREIGN KEY (grade) REFERENCES Gradepoint(grade)
```

```
ON DELETE SET NULL
```

```
-- When a grade in the Gradepoint table is deleted, it will set the grade to NULL in the  
takes table, maintaining referential integrity while avoiding orphan records.
```

```
ON UPDATE CASCADE;
```

```
-- If a grade in the Gradepoint table is updated, the corresponding records in takes will  
automatically update, ensuring consistent grade values across the database.;
```

Answer 2C

```
DROP VIEW IF EXISTS GradePointAvg;
```

```
create view GradePointAvg as
```

```
Select distinct t.id,round(sum(point*credits)/sum(credits),2) as GPA from takes t
```

```
INNER join gradepoint gp on gp.grade=t.grade
```

```
INNER join course c on c.course_id=t.course_id
```

```
where t.grade is not null
```

```
GROUP BY T.ID;
```

```
Select * from GradePointAvg;
```

3. Find the name, ID, and GPA of the Comp. Sci. student who has the highest GPA among all Comp. Sci. students. If several students are tied for the highest, your query should return them all. Do not sort students by GPA

Answer 3

```
Select ID,name,GPA from
(
  Select distinct s.id,
  S.name,
  round(sum(point*credits)/sum(credits),2) as GPA ,
  rank() over (Order by sum(point*credits)/sum(credits) DESC) as rnk
from takes t
left join gradepoint gp on gp.grade=t.grade
left join course c on c.course_id=t.course_id
left join student s on s.ID=t.ID
where t.grade is not null and s.dept_name='Comp. Sci.'
GROUP BY T.ID,s.name
) a
where rnk=1;
```

4. Find the ID and name of each Comp. Sci. student who has not taken any courses offered by the Math department.

Answer 4

```
Select distinct s.ID,s.name from Student s  
  
inner join takes t on t.ID=s.ID  
  
inner join course c on t.course_id=c.course_id  
  
where s.dept_name='Comp. Sci.' and c.dept_name not in ('Math') ;
```

;

5. Find ID, name, course_id of each student and each course they took in Fall 2009. Students who did not take any courses that semester should be listed with NULL as the course_id.

Answer 5

```
Select s.ID,name,t.course_ID from  
  
student s  
  
left join  
  
(Select t.ID,t.course_id from takes t where semester='Fall' and year='2009') t on  
t.ID=s.ID;
```

6. Find the ID, name, and Fall 2009 GPA for each student. Students who didn't take any courses in Fall 2009 should be listed with GPA either NULL or zero.

Answer 6

```
Select s.ID, name, Fall_2009_GPA from student s
left join (
  Select t.id,
  round(sum(point*credits)/sum(credits),2) as Fall_2009_GPA
  from takes t
  inner join gradepoint gp on gp.grade=t.grade
  inner join course c on c.course_id=t.course_id
  Where t.grade is not null and t.year=2009 and t.semester='Fall'
  GROUP BY T.ID) gpa on s.id=gpa.id;
```

7. Find the ID and name of each student who has taken every course taught by the instructor whose ID is 10101.

a. Write an SQL query using checks for empty set differences

b. Write an SQL query using comparison of sizes of sets

c. Write a TRC query. Hint: it should involve universal quantification.

Answer 7 a

```
SELECT s.id, s.name
FROM student s
WHERE NOT EXISTS (
  SELECT course_id
  FROM teaches
  WHERE ID = 10101
  AND (course_id, sec_id, semester, year) NOT IN (
    SELECT course_id, sec_id, semester, year
    FROM takes
    WHERE ID = s.ID
  )
);
```

Answer 7 b

```
        SELECT s.id, s.name
FROM student s
WHERE (
    SELECT COUNT(*)
    FROM teaches
    WHERE ID = 10101
) = (
    SELECT COUNT(*)
    FROM takes t2
    JOIN teaches t3 ON t2.course_id = t3.course_id
        AND t2.sec_id = t3.sec_id
        AND t2.semester = t3.semester
        AND t2.year = t3.year
    WHERE t2.ID = s.ID AND t3.ID = 10101
)
```

Answer 7 c

$$\{ t \mid \exists s \in \text{student} (t[\text{ID}] = s[\text{ID}] \wedge t[\text{name}] = s[\text{name}] \\ \wedge (\forall a \in \text{teaches} (a[\text{ID}] = 10101 \\ \exists b \in \text{takes} (b[\text{ID}] = s[\text{ID}] \wedge b[\text{semester}] = a[\text{semester}] \\ \wedge b[\text{year}] = a[\text{year}] \\ \wedge b[\text{course-id}] = a[\text{course-id}] \\ \wedge b[\text{sec-id}] = a[\text{sec-id}])))) \}$$

8. Using Retailer DB: Let's define the "profitability" of a product to be the sum of the difference between priceEach and basePrice for all items ordered (taking quantity into account) divided by the total number of items of that product ordered. Write an SQL query to find the product name and product code of the product that has the highest profitability. (If there are ties, all such products should be listed. Do not use sorting.) Your solution should include WITH, VIEW, or TEMPORARY TABLE creations with comments to make them readable.

Answer 8

```
Drop view if exists product_profit;

Create view product_profit as

select

productname,

productCode,

profitability from

-- have to put in subquery because of rank function

( select

productname,

productCode,

-- Calculating profitability : total_order_profit is derived from subquery

sum(total_order_profit)/sum(quantityOrdered) as profitability,

rank() over ( ORDER BY sum(total_order_profit)/sum(quantityOrdered) DESC)rnk

-- Used rank function to rank based on profitability

from

-- subquery to get profit per order for item which we could add later and divide by total

quantity as did above

(Select p.productname,
```

```
od.productCode,  
(od.priceEach-p.buyPrice)*quantityOrdered as total_order_profit,  
quantityOrdered  
from orderdetails od  
left join products p on p.productCode=od.productCode  
) as summary  
group by productCode) a where rnk =1  
;  
  
Select * from product_profit;
```

9. Consider the posted solution to the tennis tournament problem from HW 1. Derive schemas and CREATE TABLE definitions for relevant tables and write constraints expressing the following:

- a. A player cannot play a match against themselves**
- b. Nomore than one match can be played on a given court during a given time slot**
- c. A player cannot play in more than one match during the same time slot**

Try to determine whether your DBMS allows you to write those constraints in SQL and, if so, whether they are enforced by the DBMS. Briefly explain how you investigated this and your findings.

Answer 9.a

Drop table if exists Players;

```
CREATE TABLE Players (  
    playerID INT PRIMARY KEY,  
    name VARCHAR(100),  
    age INT,  
    ranking INT,  
    seed INT,  
    UNIQUE(seed)  
);
```

Drop table if exists Events;

```
CREATE TABLE Events (  
    eventID INT PRIMARY KEY,  
    description VARCHAR(255),  
    category ENUM('Mens Singles', 'Womens Singles', 'Juniors', 'Mixed Singles'),  
    date DATE
```

);

Drop table if exists Courts;

```
CREATE TABLE Courts (  
    courtNumber INT PRIMARY KEY,  
    surface VARCHAR(50),  
    location VARCHAR(100)  
);
```

Drop table if exists TimeSlots;

```
CREATE TABLE TimeSlots (  
    timeSlotID INT PRIMARY KEY,  
    date DATE,  
    startTime TIME,  
    endTime TIME,  
    UNIQUE(date, startTime) -- Ensures uniqueness of each time slot  
);
```

Drop table if exists SinglesMatches;

```
CREATE TABLE SinglesMatches (  
    matchID INT PRIMARY KEY,  
    player1ID INT,  
    player2ID INT,  
    eventID INT,  
    courtNumber INT,  
    timeSlotID INT,  
    score VARCHAR(20),
```

```
    highlights TEXT,  
    FOREIGN KEY (player1ID) REFERENCES Players(playerID),  
    FOREIGN KEY (player2ID) REFERENCES Players(playerID),  
    FOREIGN KEY (eventID) REFERENCES Events(eventID),  
    FOREIGN KEY (courtNumber) REFERENCES Courts(courtNumber),  
    FOREIGN KEY (timeSlotID) REFERENCES TimeSlots(timeSlotID),  
  
    CONSTRAINT chk_no_self_match CHECK (player1ID <> player2ID) -- a) Prevents a player  
    from playing against themselves  
);
```

Answer 9.b

```
ALTER TABLE SinglesMatches  
  
ADD CONSTRAINT unique_court_time UNIQUE (courtNumber, timeSlotID);
```

Answer 9.c

```
ALTER TABLE SinglesMatches  
  
ADD CONSTRAINT unique_player1_time UNIQUE (player1ID, timeSlotID),  
ADD CONSTRAINT unique_player2_time UNIQUE (player2ID, timeSlotID);
```

